


crossed run $\times$ query $\times$ judge REML

(GPT-5.2 + Sonnet, $n = 1328$)

(a) Re-running dominates the variance budget



(b) Per-run noise meets the detectable-gap floor

